

Support Vector Regression Practical Implementation - Admission dataset

Linkedin: <https://www.linkedin.com/in/satya-nerurkar-9b0655190/> (<https://www.linkedin.com/in/satya-nerurkar-9b0655190/>)

Github: <https://github.com/SatyaNerurkar> (<https://github.com/SatyaNerurkar>)

```
In [1]: 1 # Importing necessary Libraries
2 import pandas as pd
3 import numpy as np
4
5 import warnings
6 warnings.filterwarnings("ignore")
7
8 import seaborn as sns
9 import matplotlib.pyplot as plt
10
11 %matplotlib inline
```

Data Ingestion

```
In [2]: 1 raw_df = pd.read_csv("https://raw.githubusercontent.com/srinivasav22/Graduate-Admission-Prediction/master/Admission_Predict_")
```

Profile of the dataset

```
In [72]: 1 raw_df.head()
```

```
Out[72]:
```

| | Serial No. | GRE Score | TOEFL Score | University Rating | SOP | LOR | CGPA | Research | Chance of Admit |
|---|------------|-----------|-------------|-------------------|-----|-----|------|----------|-----------------|
| 0 | 1 | 337 | 118 | 4 | 4.5 | 4.5 | 9.65 | 1 | 0.92 |
| 1 | 2 | 324 | 107 | 4 | 4.0 | 4.5 | 8.87 | 1 | 0.76 |
| 2 | 3 | 316 | 104 | 3 | 3.0 | 3.5 | 8.00 | 1 | 0.72 |
| 3 | 4 | 322 | 110 | 3 | 3.5 | 2.5 | 8.67 | 1 | 0.80 |
| 4 | 5 | 314 | 103 | 2 | 2.0 | 3.0 | 8.21 | 0 | 0.65 |

```
In [73]: 1 print("This dataset contains {} rows and {} columns.".format(raw_df.shape[0],raw_df.shape[1]))
```

This dataset contains 500 rows and 9 columns.

```
In [74]: 1 raw_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Serial No.            500 non-null   int64
1   GRE Score              500 non-null   int64
2   TOEFL Score            500 non-null   int64
3   University Rating      500 non-null   int64
4   SOP                    500 non-null   float64
5   LOR                    500 non-null   float64
6   CGPA                   500 non-null   float64
7   Research               500 non-null   int64
8   Chance of Admit        500 non-null   float64
dtypes: float64(4), int64(5)
memory usage: 35.3 KB
```

```
In [75]: 1 raw_df.describe()
```

```
Out[75]:
```






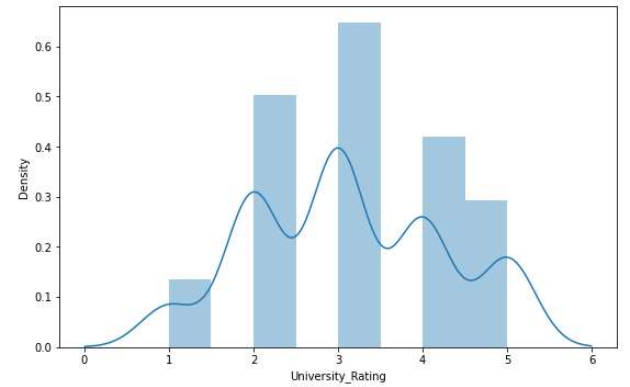



| | Serial No. | GRE Score | TOEFL Score | University Rating | SOP | LOR | CGPA | Research | Chance of Admit |
|-------|------------|------------|-------------|-------------------|------------|------------|------------|------------|-----------------|
| count | 500.000000 | 500.000000 | 500.000000 | 500.000000 | 500.000000 | 500.000000 | 500.000000 | 500.000000 | 500.000000 |
| mean | 250.500000 | 316.472000 | 107.192000 | 3.114000 | 3.374000 | 3.484000 | 8.576440 | 0.560000 | 0.72174 |
| std | 144.481833 | 11.295148 | 6.081868 | 1.143512 | 0.991004 | 0.92545 | 0.604813 | 0.496884 | 0.14114 |
| min | 1.000000 | 290.000000 | 92.000000 | 1.000000 | 1.000000 | 1.000000 | 6.800000 | 0.000000 | 0.34000 |
| 25% | 125.750000 | 308.000000 | 103.000000 | 2.000000 | 2.500000 | 3.000000 | 8.127500 | 0.000000 | 0.63000 |
| 50% | 250.500000 | 317.000000 | 107.000000 | 3.000000 | 3.500000 | 3.500000 | 8.560000 | 1.000000 | 0.72000 |
| 75% | 375.250000 | 325.000000 | 112.000000 | 4.000000 | 4.000000 | 4.000000 | 9.040000 | 1.000000 | 0.82000 |
| max | 500.000000 | 340.000000 | 120.000000 | 5.000000 | 5.000000 | 5.000000 | 9.920000 | 1.000000 | 0.97000 |

```
In [76]: raw_df.rename(columns = {'Serial No.': 'Serial_No.', 'GRE Score': 'GRE_Score', 'TOEFL Score': 'TOEFL_Score', 'University Rati
```

Exploratory Data Analysis


```
In [77]: cols = raw_df.columns
plt.figure(figsize=(20,40), facecolor='white')

for i in range(0, len(cols)):
    plt.subplot(6, 2, i+1)
    sns.distplot(x=raw_df[cols[i]],kde=True)
    plt.xlabel(cols[i])
```

```
In [78]: plt.figure(figsize=(12,6))
plt.title('Distribution of Chance of Admit feature')
sns.distplot(raw_df['Chance_of_Admit'],bins = 30,rug=True)
```

Out[78]: <AxesSubplot:title={'center':'Distribution of Chance of Admit feature'}, xlabel='Chance_of_Admit', ylabel='Density'>



```
In [3]: from pandas_profiling import ProfileReport

# EDA using pandas-profiling
profile = ProfileReport(raw_df, explorative=True)

# Displaying report in notebook cell.
profile.to_notebook_iframe()
```

Summarize dataset: 100%

72/72 [00:53<00:00, 1.62it/s, Completed]

Generate report structure: 100%

1/1 [00:15<00:00, 15.00s/it]

Render HTML: 100%

1/1 [00:08<00:00, 8.23s/it]

Overview

Dataset statistics

| | |
|-------------------------------|----------|
| Number of variables | 9 |
| Number of observations | 500 |
| Missing cells | 0 |
| Missing cells (%) | 0.0% |
| Duplicate rows | 0 |
| Duplicate rows (%) | 0.0% |
| Total size in memory | 35.3 KiB |
| Average record size in memory | 72.3 B |

Variable types

| | |
|-------------|---|
| Numeric | 7 |
| Categorical | 2 |

Alerts

| | |
|---|------------------|
| GRE Score is highly correlated with TOEFL Score and 4 other fields (TOEFL Score, University Rating, CGPA, Research, Chance of Admit.) | High correlation |
| TOEFL Score is highly correlated with GRE Score and 6 other fields (GRE Score, University Rating, SOP, LOR, CGPA, Research, Chance of Admit.) | High correlation |
| SOP is highly correlated with TOEFL Score and 4 other fields (TOEFL Score, University Rating, LOR, CGPA, Chance of Admit.) | High correlation |
| LOR is highly correlated with TOEFL Score and 3 other fields (TOEFL Score, University Rating, SOP, Chance of Admit.) | High correlation |

```
In [79]: plt.figure(figsize=(16,10))

sns.heatmap(raw_df.corr(),annot=True,cmap='icefire',linewidths=0.2)
```

Out[79]: <AxesSubplot:>



Train test split

```
In [80]: X = raw_df.drop('Chance_of_Admit', axis =1) # Independent Features
```

```
In [81]: y = raw_df["Chance_of_Admit"] # Dependent Feature
```

```
In [82]: X.head()
```

Out[82]:

| | Serial_No. | GRE_Score | TOEFL_Score | University_Rating | SOP | LOR | CGPA | Research |
|---|------------|-----------|-------------|-------------------|-----|-----|------|----------|
| 0 | 1 | 337 | 118 | 4 | 4.5 | 4.5 | 9.65 | 1 |
| 1 | 2 | 324 | 107 | 4 | 4.0 | 4.5 | 8.87 | 1 |
| 2 | 3 | 316 | 104 | 3 | 3.0 | 3.5 | 8.00 | 1 |
| 3 | 4 | 322 | 110 | 3 | 3.5 | 2.5 | 8.67 | 1 |
| 4 | 5 | 314 | 103 | 2 | 2.0 | 3.0 | 8.21 | 0 |

```
In [83]: y.head()
```

Out[83]:

| | |
|---|------|
| 0 | 0.92 |
| 1 | 0.76 |
| 2 | 0.72 |
| 3 | 0.80 |
| 4 | 0.65 |

Name: Chance_of_Admit, dtype: float64

```
In [84]: from sklearn.model_selection import train_test_split
```

```
In [85]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=42)
```



```
In [86]: print("Train dataset contains {} rows and {} columns.".format(X_train.shape[0],X_train.shape[1]))
print("Test dataset contains {} rows and {} columns.".format(X_test.shape[0],X_test.shape[1]))
print("Train dataset contains {} rows.".format(y_train.shape[0]))
print("Test dataset contains {} rows.".format(y_test.shape[0]))
```

Train dataset contains 335 rows and 8 columns.
Test dataset contains 165 rows and 8 columns.
Train dataset contains 335 rows.
Test dataset contains 165 rows.

Feature Scaling

```
In [87]: from sklearn.preprocessing import StandardScaler
```

```
In [88]: scaler = StandardScaler()
```

```
In [89]: scaler.fit(X_train)
```

```
Out[89]: StandardScaler
StandardScaler()
```

```
In [90]: print(scaler.mean_)
```

```
[255.35820896 316.7641791 107.31343284  3.09552239  3.38059701
 3.50149254  8.58498507  0.56716418]
```

```
In [91]: X_train_scaled = scaler.transform(X_train)
```

```
In [92]: X_train_scaled
```

```
Out[92]: array([[ 1.55816388e+00, -5.25174731e-01, -7.06985745e-01, ...,
                  5.38819022e-01, -8.35765678e-01, -1.14470294e+00],
                [-1.64243247e+00,  4.77035981e-01,  2.76433873e-01, ...,
                 -1.61323061e-03,  3.63045482e-01, -1.14470294e+00],
                [-1.77908715e+00, -7.98504925e-01, -1.03479228e+00, ...,
                  5.38819022e-01, -1.15657430e+00, -1.14470294e+00],
                ...,
                [ 6.73504661e-01, -1.34516531e+00, -1.36259882e+00, ...,
                 -1.62290999e+00, -2.25407747e+00, -1.14470294e+00],
                [ 1.29923923e+00, -7.07394861e-01, -3.79179206e-01, ...,
                  5.38819022e-01, -1.52803691e+00, -1.14470294e+00],
                [-1.09581377e+00, -2.51844537e-01, -2.15275936e-01, ...,
                 -1.61323061e-03, -5.65611050e-01, -1.14470294e+00]])
```

```
In [93]: X_test_scaled = scaler.transform(X_test)
```

```
In [94]: X_test_scaled
```

```
Out[94]: array([[ 7.67005229e-01,  1.57035676e+00,  1.42375676e+00, ...,
                 -1.61323061e-03,  1.61251063e+00,  8.73589088e-01],
                [-1.30439196e+00, -2.51844537e-01,  1.12530603e-01, ...,
                  5.38819022e-01,  7.68277423e-01,  8.73589088e-01],
                [ 8.60505796e-01, -1.60734472e-01, -3.79179206e-01, ...,
                 -1.08247774e+00, -1.57869090e+00, -1.14470294e+00],
                ...,
                [ 1.24889277e+00, -6.96244077e-02, -7.06985745e-01, ...,
                  1.07925128e+00,  2.61737496e-01, -1.14470294e+00],
                [-1.62085542e+00, -6.16284796e-01, -1.36259882e+00, ...,
                 -1.62290999e+00, -2.16965415e+00, -1.14470294e+00],
                [ 1.31362393e+00,  2.14856571e-02, -2.15275936e-01, ...,
                 -1.61323061e-03, -1.57869090e+00,  8.73589088e-01]])
```

Model training

```
In [95]: from sklearn.svm import SVR
```

```
In [96]: SVR_model = SVR(kernel="rbf")
```

```
In [97]: SVR_model.fit(X_train_scaled, y_train)
```

```
Out[97]: SVR
SVR()
```

```
In [98]: SVR_prediction = SVR_model.predict(X_test_scaled)
```

```
In [99]: SVR_prediction
```

```
Out[99]: array([[0.85327566, 0.70566642, 0.56577617, 0.68513912, 0.73073652,
0.84759115, 0.50884113, 0.63512354, 0.71546806, 0.75966639,
0.57183244, 0.71682881, 0.6228902 , 0.87034528, 0.80538203,
0.45539152, 0.78034293, 0.51298148, 0.52072102, 0.57781496,
0.65704711, 0.53857913, 0.7306496 , 0.78872557, 0.69357076,
0.5926933 , 0.843987 , 0.82839114, 0.59525329, 0.75285329,
0.56177445, 0.761017 , 0.57579346, 0.84035618, 0.57955442,
0.7686765 , 0.50127929, 0.88017372, 0.66595737, 0.64142824,
0.87065485, 0.52080723, 0.63350483, 0.85528765, 0.82310479,
0.47790267, 0.85620031, 0.74364715, 0.76326758, 0.85361512,
0.81739371, 0.55214794, 0.6810001 , 0.53159075, 0.90702514,
0.56836681, 0.83980413, 0.59580852, 0.6745631 , 0.45741035,
0.59644739, 0.72846957, 0.5933817 , 0.62610304, 0.48959131,
0.54214166, 0.83976541, 0.8396225 , 0.67509268, 0.69812838,
0.65151703, 0.79682174, 0.69471546, 0.58558049, 0.51274623,
0.71102862, 0.80506604, 0.81858292, 0.44022248, 0.60626074,
0.68135064, 0.78816077, 0.64312672, 0.80921569, 0.73883943,
0.65342583, 0.60539529, 0.7375616 , 0.77339084, 0.69785607,
0.74867341, 0.84240892, 0.87871142, 0.69710495, 0.78844475,
0.47504067, 0.62947709, 0.71640435, 0.73414257, 0.53127454,
0.80409686, 0.70927605, 0.52625603, 0.69155924, 0.70891608,
0.54575333, 0.74120688, 0.76668693, 0.69705305, 0.65675686,
0.61229395, 0.87815292, 0.84078298, 0.86092636, 0.40223106,
0.81305237, 0.69935842, 0.8792775 , 0.53706344, 0.60163982,
0.7232158 , 0.88432815, 0.7642912 , 0.68918653, 0.71857187,
0.72130402, 0.60175527, 0.88525868, 0.7665937 , 0.51965101,
0.5964044 , 0.66245068, 0.78301881, 0.47545906, 0.73839942,
0.41141632, 0.81003235, 0.81342825, 0.56221156, 0.71162865,
0.62552657, 0.65219413, 0.79610622, 0.47227593, 0.90804817,
0.67646231, 0.86353569, 0.76570903, 0.67480056, 0.68310625,
0.75301945, 0.59900557, 0.75347854, 0.58116682, 0.58316801,
0.67492506, 0.6141046 , 0.76912297, 0.451248 , 0.51861873,
0.62686496, 0.75779807, 0.66431766, 0.45720304, 0.61809108]])
```

Performance Metrics

```
In [100]: #Performance Metrics
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
from sklearn.metrics import r2_score
```

```
In [101]: print(mean_squared_error(y_test,SVR_prediction))
```

```
0.004878474514395206
```

```
In [102]: print(mean_absolute_error(y_test,SVR_prediction))
```

```
0.05626497201211845
```

```
In [103]: score=r2_score(y_test,SVR_prediction)
score
```

```
Out[103]: 0.7632199437959383
```

Thank you!